

CSE 3320

Operating Systems

Threads

Jia Rao

Department of Computer Science and Engineering

<http://ranger.uta.edu/~jrao>

Recap of the Last Class

- Processes
 - A program in execution
 - 5 (3)-state process model
 - Process control block
 - Execution flow
 - Resources
 - Linux processes
 - The `task_struct` structure
 - Which field stores the program counter ?
 - The `thread` field
 - Saving hardware registers can not be expressed in C, but in assembly
 - It is a processor-specific context
 - `SRC/arch/x86/include/asm/processor.h`
 - Field `ip` in struct `thread_struct`
-

Thread and Multithreading

- Process
 - resource grouping and execution
 - Thread
 - a finer-grained entity for execution and parallelism
 - Lightweight process
 - A program in execution without dedicated address space
 - Multithreading
 - Running multiple threads within a single process
-

Processes v.s. Threads

- Process

- Concurrency

- ▶ Sequential execution stream of instructions

- Protection

- ▶ A dedicated address space

- Threads

- Separate concurrency from protection

- Maintain sequential execution stream of instructions

- Share address space with other threads

A Closer Look

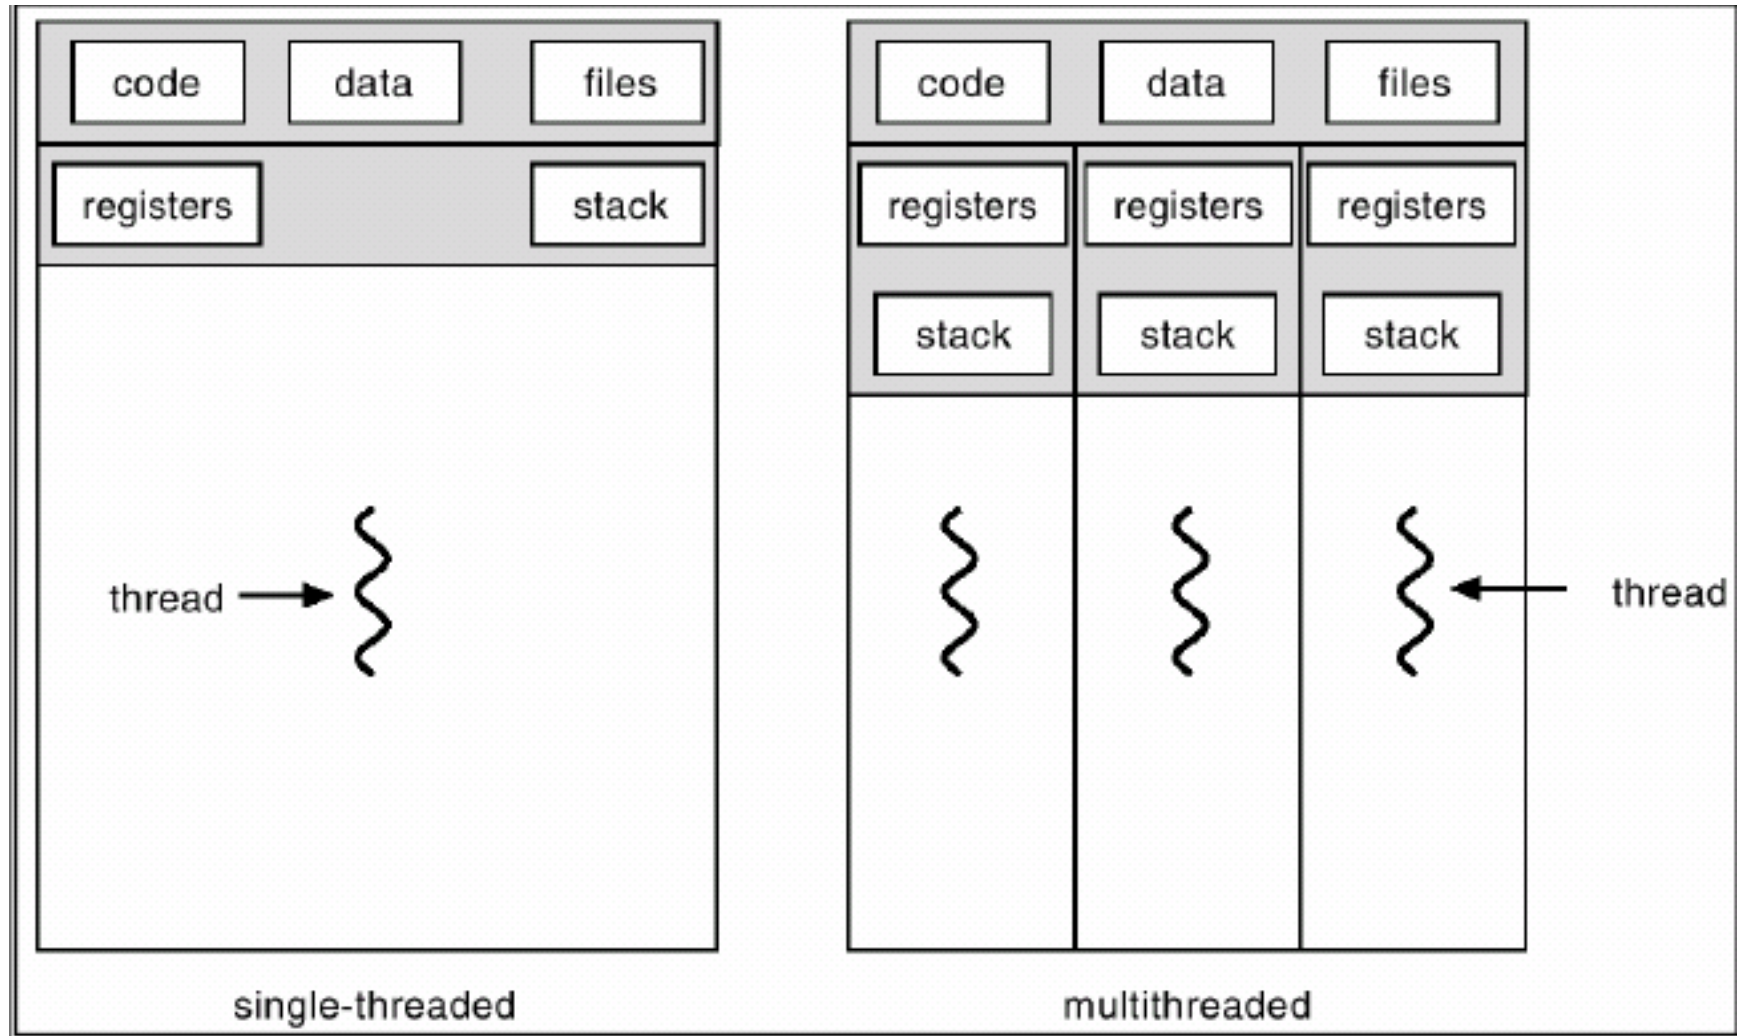
- Threads

- No data segment or heap
- Multiple can coexist in a process
- Share code, data, heap, and I/O
- Have own stack and registers
- Inexpensive to create
- Inexpensive context switching
- Efficient communication


- Processes

- Have data/code/heap
 - Include at least one thread
 - Have own address space, isolated from other processes
 - Expensive to create
 - Expensive context switching
 - IPC can be expensive
-

An Illustration



Why Multiprogramming ?

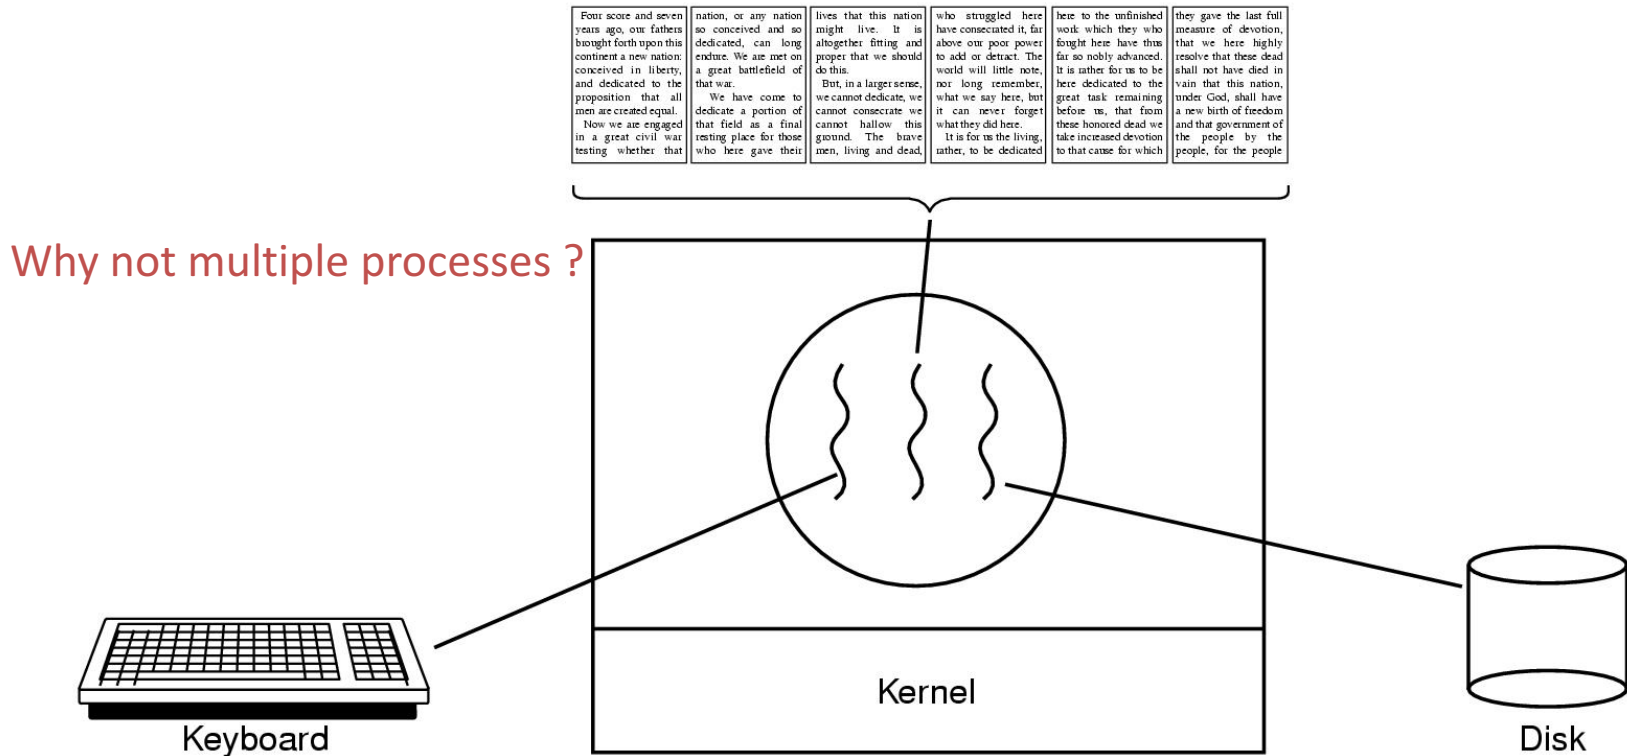


CPU utilization as a function of the number of processes in memory.

Why threads ?

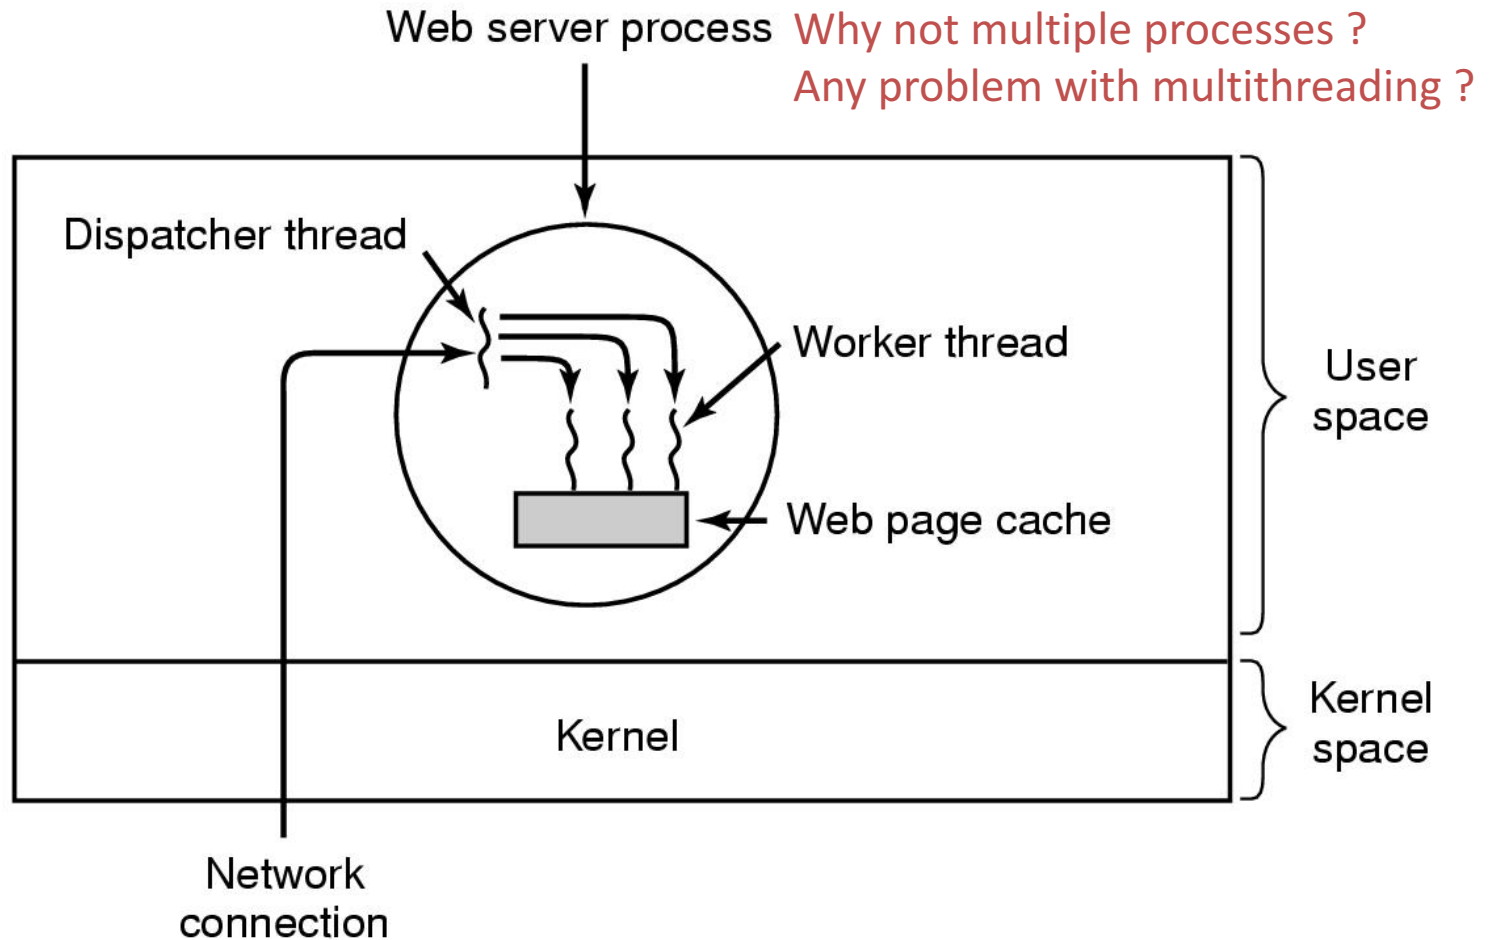
- Express concurrency
 - There are other ways to explore concurrency (e.g., non-blocking I/O), but they are difficult to program
 - Efficient communication
 - Communication can be carried out via shared data objects within the shared address space
 - Inter-process communication usually requires other OS services: file system, network system
 - Efficient creation
 - Only create the thread context
-

Thread Usage 1



A word processor with three threads.

Thread Usage 2



A multithreaded Web server.


A Simple Implementation

```
Void *worker(void *arg) // worker thread
{
    unsigned int socket;
    socket = *(unsigned int *)arg;
    process (socket);
    pthread_exit(0);
}

int main (void) // main thread, or dispatcher thread
{
    unsigned int server_s, client_s, i=0;
    pthread_t threads[200];
    server_s = socket(AF_INET, SOCK_STREAM, 0);
    .....
    listen(server_s, PEND_CONNECTIONS);
    while(1){
        client_s = accept(server_s, ...);
        pthread_create(&threads[i++], &attr, worker, &client_s);
    }
}
```

Implementing Threads in User-Space

- User-level threads: the kernel knows nothing about them



A user-level threads package

User-level Thread - Discussions

- Advantages

- No OS thread-support needed
- Lightweight: thread switching vs. process switching
 - Local procedure vs. system call (trap to kernel)
 - When we say a thread *come-to-life*? SP & PC switched
- Each process has its own customized scheduling algorithms
 - `thread_yield()`

- Disadvantages


- How blocking system calls implemented? Called by a thread?
 - Goal: to allow each thread to use blocking calls, but to prevent one blocked thread from affecting the others
 - How to change blocking system calls to non-blocking?
 - Jacket/wrapper: code to help check in advance if a call will block
 - How to deal with page faults?
 - How to stop a thread from running forever? No clock interrupts
-

Implementing Threads in the Kernel

- Kernel-level threads: when a thread blocks, kernel re-schedules another thread

- Threads known to OS
 - Scheduled by the scheduler
- Slow
 - Trap into the kernel mode
- Expensive to create and switch
 - Less expensive if in the same process
 - Registers, PC, stack pointer need to be created/changed
 - Not the memory management info

Any problem ?



A threads package managed by the kernel


Write a program that forks from a multithreaded process

Hybrid Implementations

- Use kernel-level threads and then *multiplex* user-level threads onto some or all of the kernel-level threads
- Multiplexing user-level threads Onto kernel-level threads
- Enjoy the benefits of user and kernel level threads

Too complex !
Any other problem ?

Priority inversion



Multiplexing user-level threads onto kernel-level threads

Threading Models

- N:1 (User-level threading)
 - GNU Portable Threads
 - 1:1 (Kernel-level threading)
 - Native POSIX Thread Library (NPTL)
 - M:N (Hybrid threading)
 - Solaris
-

An Example

```
void *my_thread(void *arg)
{
    int *tid = (int *)arg;

    printf("Hello from child thread: %d\n", *tid);

    return NULL;
}

int main(int argc, char *argv[]){
    pthread_t threads[NR_THREADS];

    for (i = 1; i < NR_THREADS; i++){
        tid[i] = i;

        pthread_create(&threads[i], &a, my_thread, &tid[i]);
    }

    printf("Hello from the mother thread 0 !\n");

    for (i = 1; i < NR_THREADS; ++i)
    {
        pthread_join(threads[i], NULL);
    }

    return 0;
}
```

Threads in Linux

- Thread control block (TCB)
 - The `thread_struct` structure
 - Includes registers and processor-specific context
 - Linux treats threads like processes
 - Use `clone()` to create threads instead of using `fork()`
 - `clone()` is usually not called directly but from some threading libraries, such as `pthread`.
-

Summary

- Processes v.s. threads?
 - Why threads ?
 - Concurrency + lightweight
 - Threading models
 - N:1, 1:1, M:N
 - Additional practice
 - Find out what threading model does Java belong to
 - Write (download) a simple multithreaded java program
 - When it is running, issue `ps -eLf | grep YOUR_PROG_NAME`
 - Download glibc at <http://ftp.gnu.org/gnu/glibc/> and see the nptl implementation
-