

CSE 3320

Operating Systems

CPU Scheduling

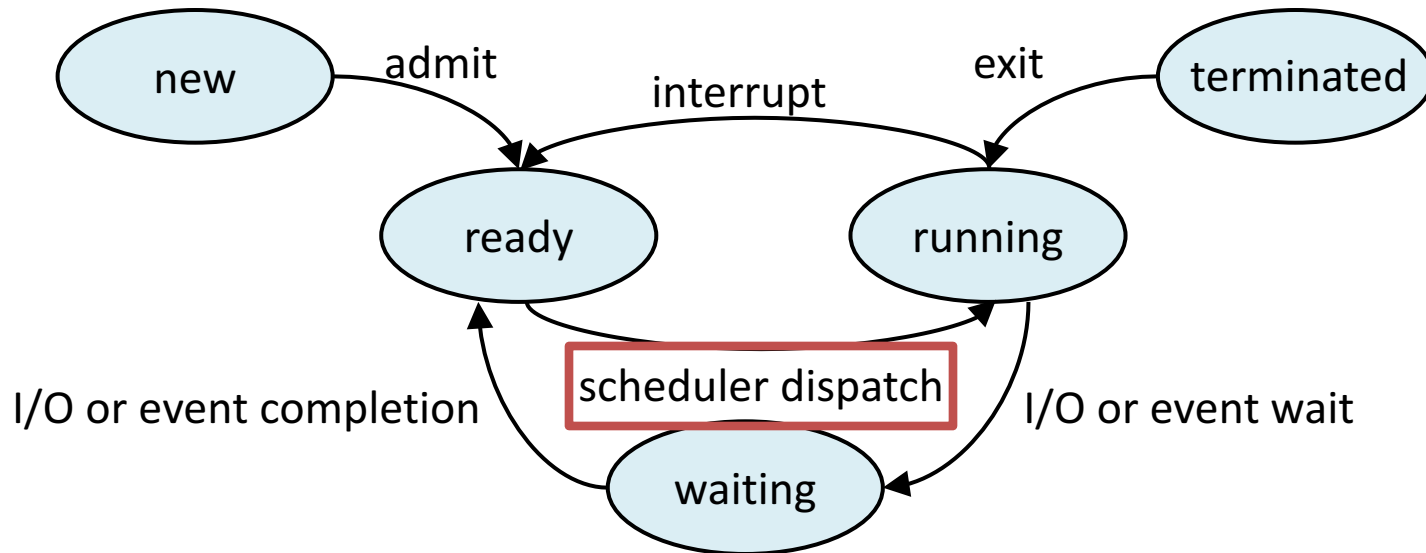
Jia Rao

Department of Computer Science and Engineering

<http://ranger.uta.edu/~jrao>

What is CPU Scheduling?

- The five-state process model



CPU scheduling

Selects from among the processes/threads that are ready to execute, and allocates the CPU to it

Why CPU Scheduling?

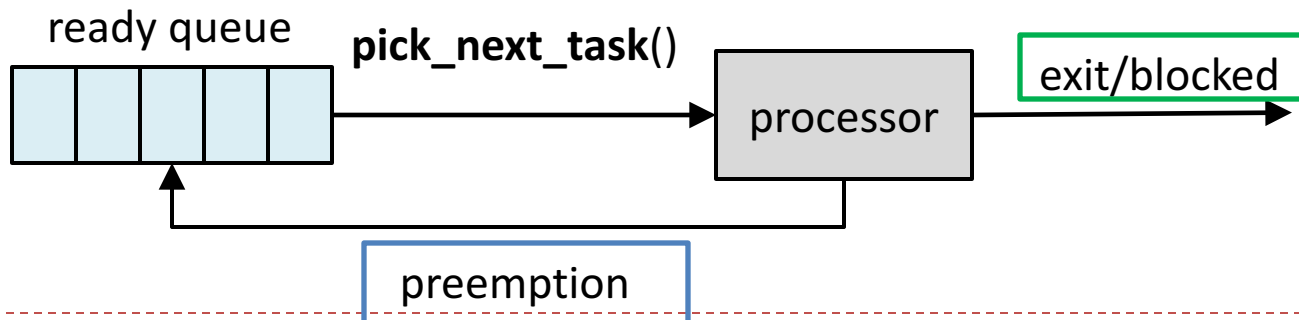
- In support of multiprogramming
 - uniprocessor systems
 - ▶ Time-sharing processor
 - multiprocessor systems
 - ▶ Efficiently distributing tasks
 - Real-time systems
 - ▶ Reliably guaranteeing deadlines
- It is (maybe) the most important part in a OS
 - Why some OS seems to be faster than others?
 - Why I do not see performance improvement when upgrading to a 16-core computer?

In this Lecture

- Outline
 - Basics of CPU scheduling
 - ▶ Scheduling policies
 - ▶ Evaluation criteria
 - ▶ Examples
 - ▶ A practical policy
 - Challenges on emerging hardware and applications
 - ▶ Many core, NUMA, asymmetric processors
 - ▶ Data center, accurate resource provisioning
 - A close look at the state-of-art
 - ▶ The Linux CFS scheduler

CPU Scheduling

- CPU scheduling may take place at
 - Clock interrupts
 - I/O completion
 - I/O interrupts
 - Termination
- Nonpreemptive
 - Scheduling only when current process terminates or gives up control
- Preemptive
 - Processes can be forced to give up control



Scheduling Goals

All systems

Fairness - giving each process a fair share of the CPU

Policy enforcement - seeing that stated policy is carried out

Balance - keeping all parts of the system busy

Batch systems

Throughput - maximize jobs per hour

Turnaround time - minimize time between submission and termination

CPU utilization - keep the CPU busy all the time

Interactive systems

Response time - respond to requests quickly

Proportionality - meet users' expectations

Real-time systems


Meeting deadlines - avoid losing data

Predictability - avoid quality degradation in multimedia systems

Scheduling Goals: A Different Point of View

- User oriented → minimize
 - Response time (wait time): the time that the first response is received (interactivity)
 - Turnaround time: the time that the task finishes
 - Predictability: variations in different runs
- System oriented → maximize
 - Throughput: # of tasks that finish per time unit
 - Utilization: the percentage of time the CPU is busy
 - Fairness: avoid starvation

Process Behaviors



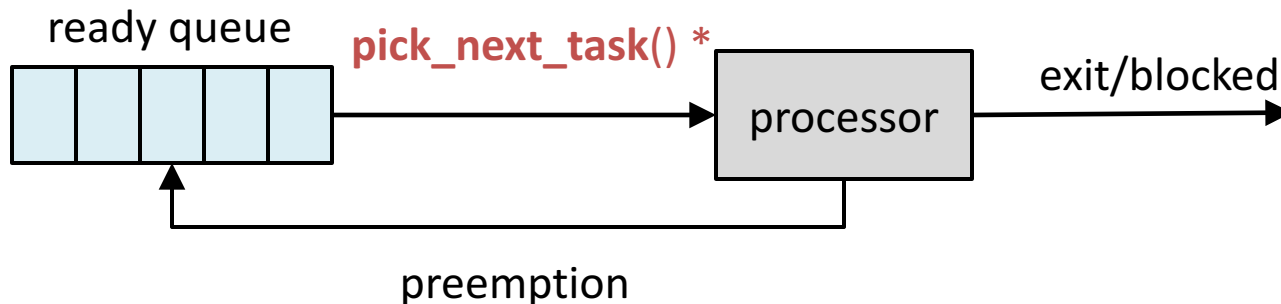
- Bursts of CPU usage alternate with periods of I/O wait
 - a **CPU-bound/CPU-intensive** process
 - an **I/O bound / I/O intensive** process
 - I/O is when a process enters the blocked state waiting for an external device to complete its work

Scheduling Policies

- Batch Systems
 - First-Come First-Serve
 - Shortest Job First
 - Shortest Remaining Time Next
- Interactive Systems
 - Round-Robin
 - Priority Scheduling
 - Multiple Queues
 - Shortest Process Next
 - Guaranteed Scheduling
 - Lottery Scheduling
- Real-time Systems
 - Rate Monotonic Scheduling
 - Earliest Deadline First Scheduling

More on Scheduling Policy



- Determine the next ready task to run
 - How we design `pick_next_task()`
- Basic policies
 - First-Come, First -Served (FCFS)
 - Shortest-Job-First (SJF)
 - Round Robin (RR)
 - Priority scheduling



First-Come, First-Serve (FCFS)

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Average turnaround time = $((8-0)+(12-1)+(26-2)+(17-3)) / 4 = 14.25$

Average response time = $(0+(8-1)+(17-2)+(12-3)) / 4 = 7.75$



Average turnaround time = $((17-0)+(5-1)+(10-3)+(26-2)) / 4 = 13$


Average response time = $(0+(1-1)+(5-3)+(17-2)) / 4 = 4.25$

Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5


$q=4$



Average turnaround time = $((20-0)+(8-1)+(26-2)+(25-3)) / 4 = 18.25$

Average response time = $(0+(4-1)+(8-2)+(12-3)) / 4 = 4.5$

$q=5$



Average turnaround time = $((22-0)+(9-1)+(26-2)+(19-3)) / 4 = 17.5$

Average response time = $(0+(5-1)+(9-2)+(14-3)) / 4 = 5.5$



Priority Scheduling

- CPU schedules the highest priority first, FCFS within the same priority

<u>Process</u>	<u>Priority</u>	<u>Burst Time</u>
P_1	3	8
P_2	1	4
P_3	4	9
P_4	2	5

P_2	P_4	P_1	P_3
-------	-------	-------	-------

<u>Process</u>	<u>Priority</u>	<u>Burst Time</u>
P_1	2	8
P_2	4	4
P_3	1	9
P_4	3	5

P_3	P_1	P_4	P_2
-------	-------	-------	-------

Put it together

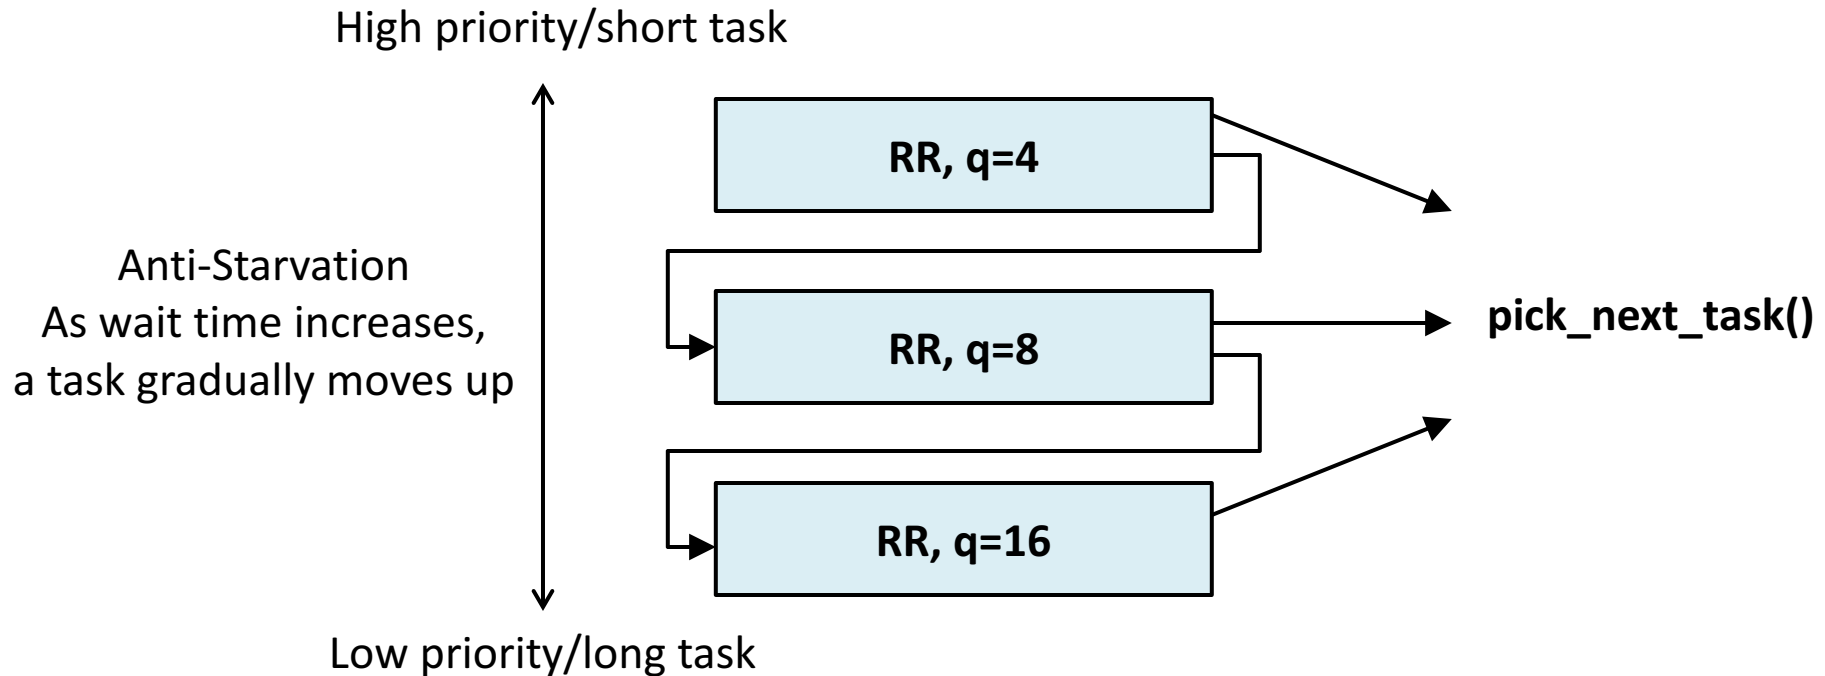
	Turnaround time	Response time
FCFS	15.25	8.75
SJF-preemptive	13	4.25
RR (q=5)	17.5	5.5
Priority scheduling	N/A	N/A

	Throughput	Response time	Starvation
FCFS	TBD	TBD	No
SJF-preemptive	High	Good	Yes
RR	Can be low	Good	No
Priority scheduling	Can be high	Can be good	Can remove



Multilevel Feedback Queue

Multilevel Feedback Queue



Windows XP, Mac OS X, Linux 2.6.22 and before

Real-time Scheduling

Schedulable real-time system


- Given
 - m periodic events
 - event i occurs within period P_i and requires C_i seconds
- Then the load can only be handled if

$$\sum_{i=1}^m \frac{C_i}{P_i} \leq 1$$

- Example: a soft real-time system with three periodic events, with periods of 100, 200, and 500 ms, respectively. If these events require 50, 30, and 100 ms of CPU time per event, respective, the system is schedulable
 - Process/context switching overhead is often an issue though!
 - Given the example, what would be the maximum CPU burst for a 4th job with a period of 500 ms ?

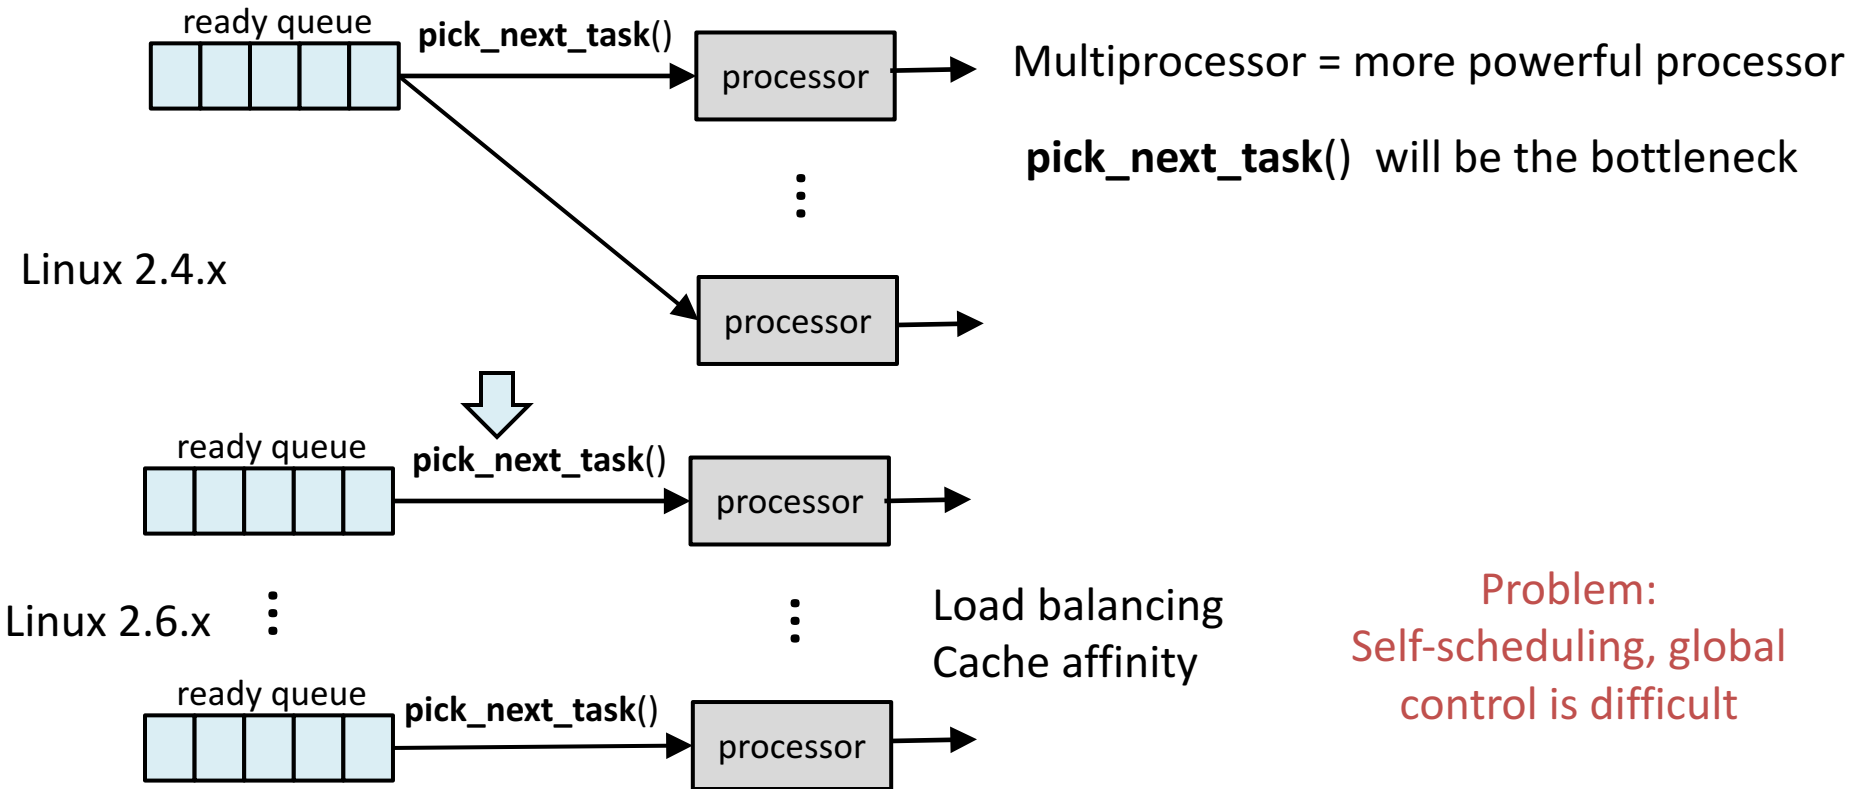
Misc.

- I/O tasks
 - Identify: runs few, sleeps a lot
 - Considered as short task
 - High priority
- Theoretical analysis
 - Assume task distribution
 - Queuing model



Challenges on Emerging Hardware and Applications

- Multiprocessor → Many core




Challenges on Emerging Hardware and Applications (cont')

- NUMA, Asymmetric processors
 - OLD: CPU time \rightarrow useful work
 - NEW: calibrated CPU time \rightarrow useful work
- Data center, accurate resource provisioning
 - Proportional fair sharing $\rightarrow P_1 : P_2 = 1 : 2$

$q=?$
When to move task?
Heuristic-based

Fine-grained
Adaptive q




A Close Look at the State-of-Art

- Linux Completely Fair Scheduler (CFS)
 - Separate ready queue per processor
 - Red-black tree based ready queue
 - Proportional fair sharing

pick_next_task()

Remove the leftmost task;
Run the task;
Every ticks update vruntime;
 $\text{vruntime} = \text{vruntime} + \text{runtime}$;
if $\text{vruntime} > \text{curr_leftmost's}$
Preempt curr task;
Put it back to the sorted tree;
Run curr_leftmost;



Fine-grain
Adaptive q
Priority: low priority
Task's vruntime runs
Faster
No starvation

Summary

- The basic scheduling policies are important
 - Multilevel Feedback Queue = RR + SJF + Priority
 - CFS = RR + SJF + Priority + smart data structure
- Additional readings
 - Go to <http://lxr.linux.no/linux+v2.6.24/kernel/>
 - Read /kernel/sched.c, /kernel/sched_fair.c, /include/linux/sched.h (starting from the schedule(void) function)
 - See how the vruntime is actually updated
 - Documentation:
<http://www.kernel.org/doc/Documentation/scheduler/sched-design-CFS.txt>
 - Another interesting scheduler: BFS
 - ▶ <http://ck.kolivas.org/patches/bfs/bfs-faq.txt>